

Harmony RPC 加速阶段性讨论

时间：2022/04/13

主题：

- Harmony 节点 RPC 测试数据
- 测试数据分析
- 给出的解决方案
- 基于负载均衡的优化
- 方案转折点
- 目前的进展

Harmony 节点 RPC 测试数据

Harmony 节点的RPC测试数据测了两版， 分别有产出文档和测试数据。

截取部分测试数据 总请求50， 并发15：

hmyv2_blockNumber

```
https://a.api.s0.t.hmny.io
hmyv2_blockNumber

Complete requests:    50
Failed requests:      0
Write errors:         0
Time per request:     1299.648 [ms] (mean)
```

hmyv2_getBlockSigners

```
https://a.api.s0.t.hmny.io
hmyv2_getBlockSigners

Complete requests:    50
Failed requests:      0
Write errors:         0
Time per request:     1217.336 [ms] (mean)
```

hmyv2_getBlocks

```
https://a.api.s0.t.hmny.io
hmyv2_getBlocks

Complete requests:    50
Failed requests:      0
Write errors:         0
Time per request:     2315.090 [ms] (mean)
```

hmyv2_getBalanceByBlockNumber

```
https://a.api.s0.t.hmny.io
hmyv2_getBalanceByBlockNumber

Complete requests:    50
Failed requests:      0
Write errors:         0
Time per request:     1394.852 [ms] (mean)
```

hmyv2_getValidatorInformation

```
https://a.api.s0.t.hmny.io
hmyv2_getValidatorInformation

Complete requests:    50
Failed requests:      0
Write errors:         0
Time per request:     1197.244 [ms] (mean)
```

hmyv2_call

```
https://a.api.s0.t.hmny.io
hmyv2_call

Complete requests:    50
Failed requests:      0
Write errors:         0
Time per request:     1187.868 [ms] (mean)
```

net_peerCount

```
https://a.api.s0.t.hmny.io
net_peerCount

Complete requests:    50
Failed requests:      0
Write errors:         0
Time per request:     1171.618 [ms] (mean)
```

hmyv2_getNodeMetadata

```
https://a.api.s0.t.hmny.io
hmyv2_getNodeMetadata

Complete requests:    50
Failed requests:      0
Write errors:         0
Time per request:     1437.167 [ms] (mean)
```

```
https://rpc.s0.t.hmny.io
hmyv2_getCurrentTransactionErrorSink

Complete requests:      50
Failed requests:        48
    (Connect: 0, Receive: 0, Length: 48, Exceptions: 0)
Write errors:           0
Non-2xx responses:      50
=====

https://rpc.s0.t.hmny.io
hmyv2_getCurrentUtilityMetrics

Complete requests:      50
Failed requests:        1
    (Connect: 0, Receive: 0, Length: 1, Exceptions: 0)
Write errors:           0
Non-2xx responses:      50
```

测试数据分析

由测试数据可以看出，从节点提供RPC服务的 **稳定性** **请求速率（RPS）** 两个维度来分析。

参考第三方团队 Infura 提供 Ethereum 的 RPC 服务，Infura 在节点服务的稳定性和请求延迟如下图：




System Metrics

Day | Week | Month


Ethereum Mainnet response time ?

249 ms



从图中可以看出，Infura提供的服务的稳定性和请求延迟上都有比较好的表现。

我们从 Harmony 官方的 Github 上也可以看出，有ISSUE也指出RPC请求过慢的问题（[Closed ISSUE Link](#)）：

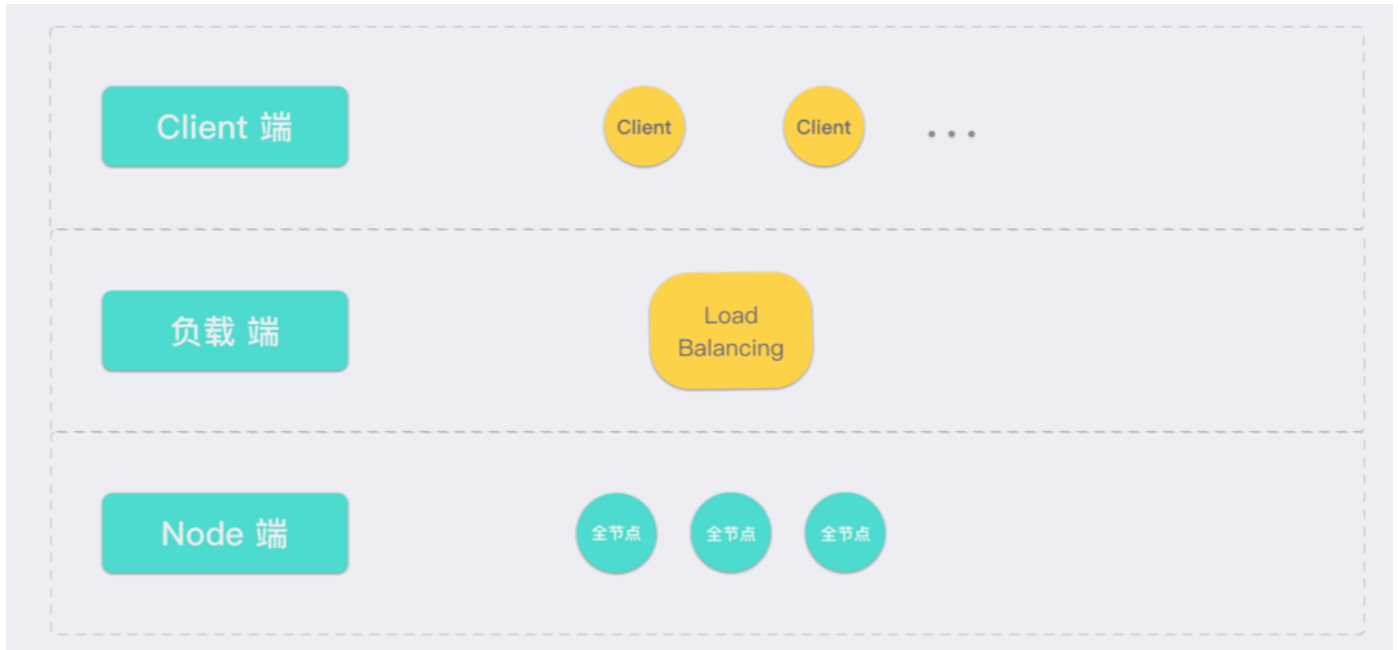


不过此ISSUE Closed，但从最近的测试结果看，还有优化的空间。

给出的解决方案

1. 堆硬件，在AWS这种云服务商上部署节点，可以考虑加大内存、使用Enhanced SSD来解决；
2. 在提供服务的内网，使用nginx做负载均衡来保证服务的可用性和较低请求延时。

我们尝试过负载均衡的策略，在内网几台性能机上部署了S0分片的节点，以Validator的模式运行（非Explorer模式）。结构图如下：



在此基础上，我们测试了服务的可用性和性能。测试数据如下：

hmyv2_blockNumber

```
https://hmyapis0.metamemo.one
hmyv2_blockNumber

Complete requests:      50
Failed requests:        0
Time per request:       148.327 [ms] (mean)
Time per request:       9.889 [ms] (mean, across all concurrent requests)
```

hmyv2_getBlockSigners

```
https://hmyapis0.metamemo.one
hmyv2_getBlockSigners

Complete requests:      50
Failed requests:        0
Time per request:       109.220 [ms] (mean)
Time per request:       7.281 [ms] (mean, across all concurrent requests)
```

hmyv2_getBlocks

```
| https://hmyapis0.metamemo.one
| hmyv2_getBlocks
|
| Complete requests:      50
| Failed requests:        0
| Time per request:       1562.746 [ms] (mean)
| Time per request:       104.183 [ms] (mean, across all concurrent requests)
```

hmyv2_getBalanceByBlockNumber

https://hmyapis0.metamemo.one
hmyv2_getBalanceByBlockNumber

Complete requests: 50
Failed requests: 0
Time per request: 140.976 [ms] (mean)
Time per request: 9.398 [ms] (mean, across all concurrent requests)

hmyv2_getValidatorInformation

https://hmyapis0.metamemo.one
hmyv2_getValidatorInformation

Complete requests: 50
Failed requests: 0
Time per request: 159.583 [ms] (mean)
Time per request: 10.639 [ms] (mean, across all concurrent requests)

hmyv2_call

https://hmyapis0.metamemo.one
hmyv2_call

Complete requests: 50
Failed requests: 0
Time per request: 134.425 [ms] (mean)
Time per request: 8.962 [ms] (mean, across all concurrent requests)

net_peerCount

https://hmyapis0.metamemo.one
net_peerCount

Complete requests: 50
Failed requests: 0
Time per request: 105.284 [ms] (mean)
Time per request: 7.019 [ms] (mean, across all concurrent requests)

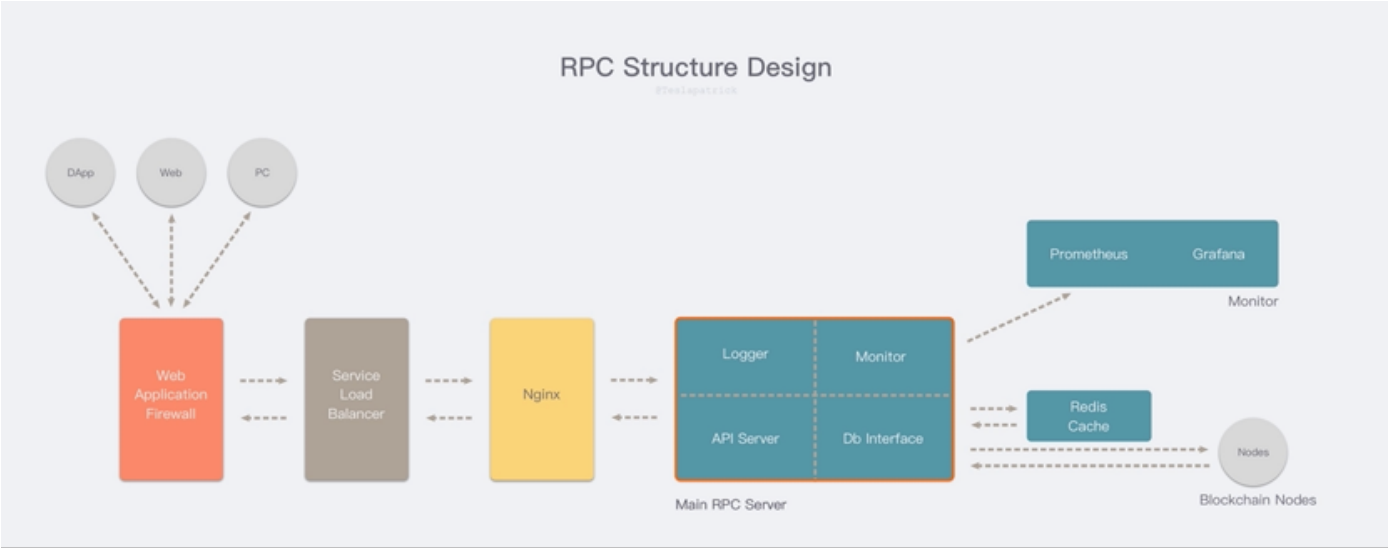
hmyv2_getNodeMetadata

https://hmyapis0.metamemo.one
hmyv2_getNodeMetadata

Complete requests: 50
Failed requests: 34
(Connect: 0, Receive: 0, Length: 34, Exceptions: 0)
Time per request: 186.314 [ms] (mean)

基于负载均衡的优化

基于讨论和问题分析的结果，针对Harmony的RPC 加速结构如下图所示。



DApp、Web、PC端的请求，经过Load Balancer，最终会到达我们的RPC Server。

其中的Db Interface主要是请求数据。数据的来源可以是Redis Cache、Nodes或是SQL Server等。由于目前的瓶颈主要在读取链上的数据，我们可以使用Redis Cache 存储一些量大且不变的数据，如果是Balance等经常更新的数据，不适合存储在Redis中，同时要考虑Redis数据的读写分离，增加Cache系统的稳定性。

Redis Cache的缓存策略，最贴近目前需求的是Cache Aside Pattern：

- Cache 未命中：DApp先从cache取链上数据，如果没有得到，则从Nodes中取数据，成功返回后放到Redis缓存中。
- Cache 命中：DApp从cache中取链上数据，取到后返回。

基于以上的结构，我们主要实现了：

1. RPC Server （RPC Proxy）的 节点Endpoint筛选 负载均衡策略开发；
 - 基于权重的负载均衡策略

```

type WeightMap struct {
    len    int
    lock   sync.RWMutex
    Nodes []*nodeWithWeight
}

type nodeWithWeight struct {
    rpc          rpc.RPCClient
    initWeight   int64
    currentWeight int64
    effectiveWeight int64
}

```

- 一致性Hash

```

type ConsistentHashMap struct {
    hash    func(data []byte) uint32
    hashMap map[uint32]rpc.RPCClient
    keys    uint32Slice
    replica int
    len     int
    lock    sync.RWMutex
}

type uint32Slice []uint32

```

2. 中转RPC的请求到节点Endpoint

```

func (l *ApiLogic) handleApiMessage(id uint64, method string, params jsonrpc.Params,
loadBalance balancing.LoadBalance) (resp *types.Response, err error) {
    // todo: 降低圈复杂度!
    switch method {
    case "getBalance":
        var address string
        params.UnmarshalSingleParam(0, &address)
        if l.useCache() {
            balance, err := l.svcCtx.Cache.GetBalance(address)
            if balance != nil {
                raw, _ := balance.MarshalJSON()

```



```

        return l.SuccessResponse(id, raw), err
    }
}
// get balance from endpoint
res, err := loadBalance.Get("getBalance").
(*hmy.HmyRpcClient).GetBalance(context.Background(), address)
if res.Error != nil {
    return l.ErrorResponse(id, res.Error), nil
}

var r *big.Int
if res.Result != nil {
    err = json.Unmarshal(*res.Result, &r)
}
// write to cache
l.svcCtx.Cache.SetBalance(address, r)

return l.SuccessResponse(id, *res.Result), nil
}

case "getBalanceByBlockNumber":
    ...
    ...
    ...
}

```

3. Redis 缓存链上数据的功能开发

```

type Config struct {
    Endpoint string `json:"endpoint"`
    Password string `json:"password"`
    Database int    `json:"database"`
    PoolSize int    `json:"poolSize"`
}

type RedisClient struct {
    client *redis.Client
    prefix string
    timeout time.Duration
}

func (r *RedisClient) SetBalance(address string, balance *big.Int) error {

```

```

pipe := r.client.TxPipeline()
defer pipe.Close()
// set balance
pipe.HMSet(context.Background(), r.formatKey("balanceLatest", address), map[string]interface{}{
    "latest": balance.String(),
})
// set expire time
pipe.Expire(context.Background(), r.formatKey("balanceLatest", address), r.timeout)
_, err := pipe.Exec(context.Background())
return err
}

```

方案转折点

在测试代码时用官方的 `net/http` 包替换为 `fasthttp`，在第一次连接时建立连接比较久（不同的endpoint不同时间），`fasthttp`在 `DefaultMaxIdleConnDuration = 10 * time.Second` 内不会释放TCP连接，这样在 `DefaultMaxIdleConnDuration` 时间内，可以再次发起请求，而不需要重新建立连接。

这样的连接的请求情况如图：

图。

同时`fasthttp`默认的连接数为 `DefaultMaxConnsPerHost = 512`，这样在并发测试下，会有更好的表现。

基于`fasthttp`测试的情况，可以把请求后端endpoint的连接统一用 `连接池` 替代 `net/http`。因为经过`nginx`负载均衡后，RPC Server（RPC Proxy）跟后端Endpoint之间不需要频繁建立 / 断开 TCP连接。

这样从RPC请求过来，经过`nginx`、RPC Server的负载均衡，再到Endpoint的可以如下图：

图。

目前的进展

目前在RPC加速主要的开发如下：

1. Endpoint端的负载均衡；
2. 连接池；
3. Redis Cache 缓存无状态的链上数据。

基于开发1、2功能的基础上的测试数据：

hmyv2_blockNumber

```
http://103.39.231.220:8888/api/v1/hmy/v2
hmyv2_blockNumber

Complete requests:    50
Failed requests:      0
Time per request:     244.884 [ms] (mean)
Time per request:     16.326 [ms] (mean, across all concurrent requests)
```

hmyv2_getBlockSigners

```
http://103.39.231.220:8888/api/v1/hmy/v2
hmyv2_getBlockSigners

Complete requests:    50
Failed requests:      0
Time per request:     265.416 [ms] (mean)
Time per request:     17.694 [ms] (mean, across all concurrent requests)
```

hmyv2_getBlocks

```
http://103.39.231.220:8888/api/v1/hmy/v2
hmyv2_getBlocks

Complete requests:    50
Failed requests:      0
Time per request:     1611.972 [ms] (mean)
Time per request:     107.465 [ms] (mean, across all concurrent requests)
```

hmyv2_getBalanceByBlockNumber

```
http://103.39.231.220:8888/api/v1/hmy/v2
hmyv2_getBalanceByBlockNumber

Complete requests:    50
Failed requests:      0
Time per request:     105.700 [ms] (mean)
Time per request:      7.047 [ms] (mean, across all concurrent requests)
```

hmyv2_getValidatorInformation

```
http://103.39.231.220:8888/api/v1/hmy/v2
hmyv2_getValidatorInformation

Complete requests:    50
Failed requests:      0
Time per request:     108.739 [ms] (mean)
Time per request:      7.249 [ms] (mean, across all concurrent requests)
```

hmyv2_call

```
http://103.39.231.220:8888/api/v1/hmy/v2
hmyv2_call
```

```
Complete requests:    50
Failed requests:      0
Time per request:     119.941 [ms] (mean)
Time per request:     7.996 [ms] (mean, across all concurrent requests)
```

net_peerCount

```
http://103.39.231.220:8888/api/v1/hmy/v2
net_peerCount
```

```
Complete requests:    50
Failed requests:      0
Time per request:     88.574 [ms] (mean)
Time per request:     5.905 [ms] (mean, across all concurrent requests)
```

hmyv2_getNodeMetadata

```
http://103.39.231.220:8888/api/v1/hmy/v2
hmyv2_getNodeMetadata
```

```
Complete requests:    50
Failed requests:      0
Time per request:     85.439 [ms] (mean)
Time per request:     5.696 [ms] (mean, across all concurrent requests)
```

基于Redis Cache的测试数据:

待补充。